

# Churn Prediction Model

June 26, 2024

## 1 Importing Necessary Libraries

```
[19]: import pandas as pd #Data Manipulation
import numpy as np
import seaborn as sns #Data Visualization
import matplotlib.pyplot as plt #Data Visualization
import plotly.express as px #Data Visualization
import plotly.graph_objects as go #Data Visualization
```

## 2 Load the Dataset

```
[20]: df = pd.read_csv("Churn_ Data.csv")
```

## 3 Understanding the Dataset

```
[21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25000 entries, 0 to 24999
Columns: 111 entries, s6.new.rev.p2.m2 to s3.rev.p1
dtypes: float64(80), int64(31)
memory usage: 21.2 MB
```

```
[22]: df.head()
```

```
[22]:  s6.new.rev.p2.m2  s1.new.rev.m1  s3.og.rev.4db.p5  s3.new.rev.4db.p5  \
0             -0.76         88.0482         3.106604         3.754955
1             -0.98         67.5039         3.094574         5.550865
2             -0.98         33.9248         2.324016         2.438114
3             -0.92         82.6780         2.630749         2.858961
4             -0.97         96.8379         2.674316         2.912397

      s4.usg.ins.p2  s4.og.unq.any.p2  s2.rch.val.p6  s1.og.rev.all.m1  \
0                4                14         39.29         57.320
1                1                 2         21.67         38.700
```

2	2	3	30.00	15.320
3	2	3	50.00	51.956
4	3	2	22.50	66.886

	s8.new.rev.p6	s4.loc.ic.ins.p1	...	prop.og.mou.tot.mou.all.p6	\
0	-0.17	1	...	0.454642	
1	-0.32	3	...	0.343190	
2	-0.05	3	...	0.101838	
3	-0.18	4	...	0.066602	
4	0.01	4	...	0.219821	

	prop.i2i.og.mou.p6	s4.loc.ic.ins.p2	s4.std.ic.ins.l14	\
0	0.497397	4	0	
1	0.767617	6	0	
2	0.619034	6	1	
3	0.437088	7	2	
4	0.585977	6	1	

	s4.low.blnc.ins.p4	s3.og.rev.all.m2	s3.new.rev.m2	prop.og.mou.any.p6	\
0	9	6.02	8.20	46.465636	
1	20	3.66	8.10	34.525456	
2	19	4.33	4.36	10.298451	
3	11	3.40	3.53	6.670783	
4	14	3.85	3.87	21.998905	

	prop.loc.i2i.mou.og.mou.p3	s3.rev.p1
0	0.609456	0.22
1	1.000000	0.38
2	0.699592	0.11
3	0.086617	5.18
4	0.683105	0.10

[5 rows x 111 columns]

[23]: df.tail()

[23]:	s6.new.rev.p2.m2	s1.new.rev.m1	s3.og.rev.4db.p5	s3.new.rev.4db.p5	\
24995	0.21	132.0365	2.652236	2.857739	
24996	0.80	77.0154	3.763389	5.012503	
24997	0.01	148.8337	3.823940	4.334250	
24998	0.17	1012.4398	14.667580	14.579567	
24999	-1.00	275.3530	5.134579	5.954062	

	s4.usg.ins.p2	s4.og.unq.any.p2	s2.rch.val.p6	s1.og.rev.all.m1	\
24995	5	8	26.67	123.396	
24996	2	8	27.88	62.140	
24997	6	10	10.00	98.900	

24998	7	67	42.92	734.005
24999	1	1	53.50	250.340

	s8.new.rev.p6	s4.loc.ic.ins.p1	...	prop.og.mou.tot.mou.all.p6	\
24995	-0.16	4	...	0.145831	
24996	0.19	4	...	0.529829	
24997	-0.03	2	...	0.327245	
24998	0.70	4	...	0.824671	
24999	-0.48	4	...	0.377281	

	prop.i2i.og.mou.p6	s4.loc.ic.ins.p2	s4.std.ic.ins.l14	\
24995	0.200151	7	0	
24996	0.169835	7	0	
24997	0.407944	3	0	
24998	0.889239	7	1	
24999	0.609046	7	0	

	s4.low.blnc.ins.p4	s3.og.rev.all.m2	s3.new.rev.m2	\
24995	18	3.57	3.83	
24996	18	6.89	7.70	
24997	12	6.63	7.48	
24998	1	19.36	22.26	
24999	18	5.42	8.02	

	prop.og.mou.any.p6	prop.loc.i2i.mou.og.mou.p3	s3.rev.p1
24995	14.896154	0.328027	0.76
24996	55.156230	0.288006	12.74
24997	33.222018	0.235918	8.07
24998	82.549378	0.952962	21.21
24999	38.590040	1.000000	0.00

[5 rows x 111 columns]

```
[24]: print(df.shape)
```

(25000, 111)

```
[25]: df.describe().T
```

	count	mean	std	min	\
s6.new.rev.p2.m2	25000.0	-0.003730	2.727916	-1.000000	
s1.new.rev.m1	25000.0	281.073083	276.075983	0.000000	
s3.og.rev.4db.p5	25000.0	4.890003	4.212452	0.000000	
s3.new.rev.4db.p5	25000.0	7.070194	6.318992	0.000833	
s4.usg.ins.p2	25000.0	5.460080	2.184444	0.000000	
...	...	...	...	...	
s3.og.rev.all.m2	25000.0	8.008660	6.152429	0.000000	

s3.new.rev.m2	25000.0	12.540182	11.540611	0.000000
prop.og.mou.any.p6	25000.0	53.594165	21.408486	0.000000
prop.loc.i2i.mou.og.mou.p3	25000.0	0.483975	0.292349	0.000000
s3.rev.p1	25000.0	9.951366	17.648128	0.000000
	25%	50%	75%	max
s6.new.rev.p2.m2	-0.580000	-0.170000	0.280000	316.860000
s1.new.rev.m1	101.563800	204.859600	370.711650	5702.924300
s3.og.rev.4db.p5	2.367288	3.729944	5.993342	153.221695
s3.new.rev.4db.p5	3.318825	5.231268	8.395736	170.200441
s4.usg.ins.p2	5.000000	7.000000	7.000000	7.000000
...	...	...	...	...
s3.og.rev.all.m2	4.207500	6.345000	9.830000	171.780000
s3.new.rev.m2	6.167500	9.350000	14.620000	386.480000
prop.og.mou.any.p6	39.378142	53.976203	68.312416	100.000000
prop.loc.i2i.mou.og.mou.p3	0.251304	0.477621	0.716538	1.000000
s3.rev.p1	1.970000	5.380000	11.400000	585.500000

[111 rows x 8 columns]

## 4 Processing of the Data

### 4.0.1 Check for Misclassified Data Types

```
[26]: misclassified_columns = []
      for col in df.columns:
          if df[col].dtype == 'object':
              try:
                  df[col] = pd.to_numeric(df[col])
              except ValueError:
                  misclassified_columns.append(col)
      print("Misclassified columns:", misclassified_columns)
```

Misclassified columns: []

### 4.0.2 Check for NULL Values

```
[27]: df.isnull().sum()
      df.isnull().sum() / df.shape[0] * 100
```

```
[27]: s6.new.rev.p2.m2      0.0
      s1.new.rev.m1        0.0
      s3.og.rev.4db.p5     0.0
      s3.new.rev.4db.p5    0.0
      s4.usg.ins.p2        0.0
      ...
      s3.og.rev.all.m2     0.0
```

```
s3.new.rev.m2          0.0
prop.og.mou.any.p6     0.0
prop.loc.i2i.mou.og.mou.p3  0.0
s3.rev.p1              0.0
Length: 111, dtype: float64
```

#### 4.0.3 Check for Duplicate Values

```
[28]: duplicates = df.duplicated().sum()
df = df.drop_duplicates()
print(f"Removed {duplicates} duplicate rows")
print("New shape:", df.shape)
```

```
Removed 0 duplicate rows
New shape: (25000, 111)
```

#### 4.0.4 Check for Unique values

```
[29]: unique_value_columns = [col for col in df.columns if df[col].nunique() == df.
    ↪shape[0]]
df = df.drop(columns=unique_value_columns)
print("Removed unique value columns:", unique_value_columns)
print("New shape:", df.shape)
```

```
Removed unique value columns: []
New shape: (25000, 111)
```

#### 4.0.5 Check for Zero Variance variables and Removal

```
[30]: zero_variance_columns = [col for col in df.columns if df[col].std() == 0]
df = df.drop(columns=zero_variance_columns)
print("Removed zero variance columns:", zero_variance_columns)
print("New shape:", df.shape)
```

```
Removed zero variance columns: []
New shape: (25000, 111)
```

#### 4.0.6 Outliers treatment (using IQR method and +/- Sigma Approach)

```
[31]: # Step 4.6: Outliers treatment (using IQR method and +/- 3 Sigma Approach)
def treat_outliers(df):
    for column in df.select_dtypes(include=[np.number]).columns:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1

    # IQR method
```

```

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# +/- 3 Sigma approach (assuming normal distribution)
mean = df[column].mean()
std = df[column].std()
sigma_lower_bound = mean - 3 * std
sigma_upper_bound = mean + 3 * std

# Combining IQR and sigma methods (consider the more extreme bound)
final_lower_bound = max(lower_bound, sigma_lower_bound)
final_upper_bound = min(upper_bound, sigma_upper_bound)

df[column] = np.clip(df[column], final_lower_bound, final_upper_bound)
return df

df = treat_outliers(df)
print(df.describe())
print("New shape:", df.shape)

```

	s6.new.rev.p2.m2	s1.new.rev.m1	s3.og.rev.4db.p5	s3.new.rev.4db.p5 \
count	25000.000000	25000.000000	25000.000000	25000.000000
mean	-0.083814	263.414113	4.559075	6.492778
std	0.668743	211.173061	2.897105	4.219776
min	-1.000000	0.000000	0.000000	0.000833
25%	-0.580000	101.563800	2.367288	3.318825
50%	-0.170000	204.859600	3.729944	5.231268
75%	0.280000	370.711650	5.993342	8.395736
max	1.570000	774.433425	11.432423	16.011101

	s4.usg.ins.p2	s4.og.unq.any.p2	s2.rch.val.p6	s1.og.rev.all.m1 \
count	25000.000000	25000.000000	25000.000000	25000.000000
mean	5.624880	27.137520	67.150729	201.146059
std	1.818213	23.163585	46.538585	167.480617
min	2.000000	0.000000	0.000000	0.000000
25%	5.000000	9.000000	33.000000	74.420000
50%	7.000000	21.000000	52.260000	151.168500
75%	7.000000	39.000000	89.852500	284.265000
max	7.000000	84.000000	175.131250	599.032500

	s8.new.rev.p6	s4.loc.ic.ins.p1	... prop.og.mou.tot.mou.all.p6 \
count	25000.000000	25000.000000	25000.000000
mean	-0.026442	3.374020	0.538407
std	0.253097	0.915725	0.209203
min	-0.565000	1.500000	0.000000
25%	-0.160000	3.000000	0.394227
50%	-0.020000	4.000000	0.539354

75%	0.110000	4.000000	...	0.682695
max	0.515000	4.000000	...	1.000000

	prop.i2i.og.mou.p6	s4.loc.ic.ins.p2	s4.std.ic.ins.l14	\
count	25000.000000	25000.000000	25000.000000	
mean	0.485523	5.835480	1.322600	
std	0.271146	1.667562	1.822707	
min	0.000000	2.000000	0.000000	
25%	0.274034	5.000000	0.000000	
50%	0.476759	7.000000	0.000000	
75%	0.694104	7.000000	2.000000	
max	1.000000	7.000000	5.000000	

	s4.low.blnc.ins.p4	s3.og.rev.all.m2	s3.new.rev.m2	\
count	25000.000000	25000.000000	25000.000000	
mean	8.382160	7.578758	11.333164	
std	8.961016	4.477694	6.941543	
min	0.000000	0.000000	0.000000	
25%	1.000000	4.207500	6.167500	
50%	5.000000	6.345000	9.350000	
75%	14.000000	9.830000	14.620000	
max	30.000000	18.263750	27.298750	

	prop.og.mou.any.p6	prop.loc.i2i.mou.og.mou.p3	s3.rev.p1
count	25000.000000	25000.000000	25000.000000
mean	53.594165	0.483975	7.817712
std	21.408486	0.292349	7.555598
min	0.000000	0.000000	0.000000
25%	39.378142	0.251304	1.970000
50%	53.976203	0.477621	5.380000
75%	68.312416	0.716538	11.400000
max	100.000000	1.000000	25.545000

[8 rows x 111 columns]  
New shape: (25000, 111)

#### 4.0.7 Handling Missing Values

```
[32]: def treat_missing_values(df):
    missing_percentage = df.isnull().mean() * 100
    columns_to_drop = missing_percentage[missing_percentage > 50].index
    df = df.drop(columns=columns_to_drop)
    print(f"Dropped columns with >50% missing values: {list(columns_to_drop)}")

    total_missing_percentage = df.isnull().sum().sum() / (df.shape[0] * df.
↪shape[1]) * 100
    if total_missing_percentage < 5:
```

```

        df = df.dropna()
        print("Removed records with missing values (less than 5% of total)")
    else:
        for column in df.columns:
            if df[column].dtype in ['int64', 'float64']:
                df[column].fillna(df[column].median(), inplace=True)
            elif df[column].dtype == 'object':
                df[column].fillna(df[column].mode()[0], inplace=True)
        return df

df = treat_missing_values(df)
print("New shape after handling missing values:", df.shape)

```

Dropped columns with >50% missing values: []  
 Removed records with missing values (less than 5% of total)  
 New shape after handling missing values: (25000, 111)

#### 4.0.8 Removing Highly Correlated Variables

```

[33]: def remove_highly_correlated(df, threshold=0.9):
        corr_matrix = df.corr().abs()
        upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).
        ↪astype(bool))
        to_drop = [column for column in upper.columns if any(upper[column] >
        ↪threshold)]
        df = df.drop(columns=to_drop)
        print(f"Removed highly correlated columns: {to_drop}")
        return df

df = remove_highly_correlated(df)
print("New shape after removing highly correlated variables:", df.shape)

```

Removed highly correlated columns: ['s1.og.rev.all.m1', 's2.rch.val.l67', 's7.new.rev.p2.p6', 's7.rtd.mou.p2.p6', 's1.new.rev.p1', 's1.rtd.mou.p1', 's1.og.rev.all.p1', 's1.og.mou.all.p1', 'snd.dec.p2', 's1.og.mou.all.p2', 's8.og.rev.p6', 's1.og.hom.mou.p2', 's5.og.rev.all.p1', 's1.og.rev.all.p2', 's1.rtd.mou.p2', 's5.rtd.mou.p1', 's1.og.mou.any.p2', 's1.hom.rmg.rev.p2', 's5.og.mou.all.p1', 's5.og.hom.mou.p1', 's4.usg.ins.p1', 's2.s4.day.no.mou.p2', 's5.s4.day.no.mou.p2', 'tot.s4.day.no.mou.p2', 's1.rev.p1', 's4.og.any.p2', 's1.loc.og.mou.p2', 's5.new.rev.p2', 's5.new.rev.p1', 's4.low.blnc.ins.l14', 's3.og.hom.mou.p1', 's3.new.rev.p2', 'tot.s4.day.no.mou.p3', 's5.og.mou.all.p2', 's4.usg.ins.l14', 's4.loc.og.ins.p2', 's3.rtd.mou.p1', 's7.s5.s4.day.nomou.p2', 's5.og.hom.mou.p2', 'prop.og.mou.tot.mou.all.p2', 's7.s5.s4.day.nomou.p3', 's3.og.rev.3db.p5', 's8.rtd.mou.p6', 's4.low.blnc.ins.p2', 's4.low.blnc.ins.m2', 's4.dec.ins.p2', 's1.rev.p2', 'prop.i2i.og.mou.p6', 's4.loc.ic.ins.p2', 's4.std.ic.ins.l14', 's4.low.blnc.ins.p4', 's3.og.rev.all.m2', 's3.new.rev.m2', 'prop.og.mou.any.p6', 's3.rev.p1']



New shape after removing highly correlated variables: (25000, 56)

#### 4.0.9 Multicollinearity (VIF > 5)

```
[34]: import warnings
warnings.filterwarnings("ignore")
from statsmodels.stats.outliers_influence import variance_inflation_factor

def calculate_vif(df):
    vif_data = pd.DataFrame()
    vif_data["feature"] = df.columns
    vif_data["VIF"] = [variance_inflation_factor(df.values, i) for i in
↳range(len(df.columns))]
    return vif_data

def remove_high_vif_features(df, threshold=5):
    while True:
        vif_data = calculate_vif(df)
        if vif_data['VIF'].max() <= threshold:
            break
        feature_to_remove = vif_data.loc[vif_data['VIF'].idxmax(), 'feature']
        df = df.drop(columns=[feature_to_remove])
        print(f"Removed feature with high VIF: {feature_to_remove}")
    return df

df = remove_high_vif_features(df)
print("Final shape after removing features with high VIF:", df.shape)

# Step 5: Final preprocessed dataset
print(df.describe())
print("Final dataset shape:", df.shape)
```

```
Removed feature with high VIF: s4.loc.ins.l14
Removed feature with high VIF: s7.rtd.mou.l21.p6
Removed feature with high VIF: s4.loc.ic.ins.l14
Removed feature with high VIF: s4.usg.ins.p2
Removed feature with high VIF: s7.new.rev.p3.p6
Removed feature with high VIF: s4.loc.og.ins.l14
Removed feature with high VIF: s1.new.rev.p2
Removed feature with high VIF: s7.new.rev.l21.p6
Removed feature with high VIF: s1.new.rev.m2
Removed feature with high VIF: s3.og.mou.all.p1
Removed feature with high VIF: s7.rtd.mou.p3.p6
Removed feature with high VIF: s4.loc.ic.ins.p1
Removed feature with high VIF: s1.og.hom.mou.p1
Removed feature with high VIF: prop.og.mou.tot.mou.all.p6
Removed feature with high VIF: s3.new.rev.4db.p5
Removed feature with high VIF: s7.rev.p2.p6
```

Removed feature with high VIF: prop.loc.i2i.mou.og.mou.p6

Removed feature with high VIF: s4.loc.og.ins.p1

Removed feature with high VIF: s1.og.hom.rev.p2

Removed feature with high VIF: s5.rev.p1

Removed feature with high VIF: s4.low.blnc.ins.p6

Removed feature with high VIF: s3.new.rev.p3

Removed feature with high VIF: s3.og.rev.all.p1

Removed feature with high VIF: s1.new.rev.m1

Removed feature with high VIF: s4.dec.ins.l14

Removed feature with high VIF: s5.rev.p2

Removed feature with high VIF: s3.og.mou.all.p2

Removed feature with high VIF: s4.og.unq.any.p2

Removed feature with high VIF: s3.og.rev.4db.p5

Removed feature with high VIF: s8.og.rev.p3

Removed feature with high VIF: s6.new.rev.p2.m2

Removed feature with high VIF: s2.s4.day.no.mou.p3

Removed feature with high VIF: prop.og.mou.any.p2

Final shape after removing features with high VIF: (25000, 23)

	s2.rch.val.p6	s8.new.rev.p6	s8.mbl.p2	s7.s4.day.no.mou.p2.p4 \
count	25000.000000	25000.000000	25000.000000	25000.000000
mean	67.150729	-0.026442	-0.547296	28.481545
std	46.538585	0.253097	3.921211	44.601940
min	0.000000	-0.565000	-7.960000	0.000000
25%	33.000000	-0.160000	-2.560000	0.100000
50%	52.260000	-0.020000	-0.080000	0.363636
75%	89.852500	0.110000	1.040000	99.000000
max	175.131250	0.515000	6.440000	99.000000

	s7.s5.s4.day.nomou.p4	s8.ic.mou.all.p3	target \
count	25000.000000	25000.000000	25000.000000
mean	0.284181	-0.038190	0.316680
std	0.333501	0.480536	0.465191
min	0.000000	-1.040000	0.000000
25%	0.000000	-0.290000	0.000000
50%	0.000000	-0.020000	0.000000
75%	0.500000	0.210000	1.000000
max	1.000000	0.960000	1.000000

	s6.rtd.mou.p2.m2	ds.usg.p6	ds.og.usg.p4	...	s1.loc.og.mou.p1 \
count	25000.000000	25000.0	25000.0	...	25000.000000
mean	-0.080950	0.0	0.0	...	30.365185
std	0.646247	0.0	0.0	...	32.871037
min	-1.000000	0.0	0.0	...	0.000000
25%	-0.540000	0.0	0.0	...	4.483200
50%	-0.150000	0.0	0.0	...	17.999900
75%	0.270000	0.0	0.0	...	45.499325
max	1.485000	0.0	0.0	...	107.023513

	s4.low.blnc.ins.p3	s8.rtd.mou.p3	s8.og.mou.all.p6	s7.rtd.mou.m1.m2	\
count	25000.000000	25000.000000	25000.000000	25000.000000	
mean	4.335200	-0.087113	-0.019145	0.897349	
std	5.027382	0.633841	0.241952	0.520586	
min	0.000000	-1.345000	-0.500000	0.000000	
25%	0.000000	-0.400000	-0.140000	0.552297	
50%	2.000000	-0.050000	-0.020000	0.860550	
75%	7.000000	0.230000	0.100000	1.176028	
max	15.000000	1.175000	0.460000	2.111626	

	s8.rev.p6	s4.rch.val.gt.30.p2	s4.std.ins.l14	s4.data.ins.l14	\
count	25000.000000	25000.000000	25000.000000	25000.000000	
mean	-0.017048	0.701160	2.103420	1.243200	
std	0.254273	0.802347	2.712373	1.965291	
min	-0.555000	0.000000	0.000000	0.000000	
25%	-0.150000	0.000000	0.000000	0.000000	
50%	-0.010000	1.000000	1.000000	0.000000	
75%	0.120000	1.000000	3.000000	2.000000	
max	0.525000	2.500000	7.500000	5.000000	

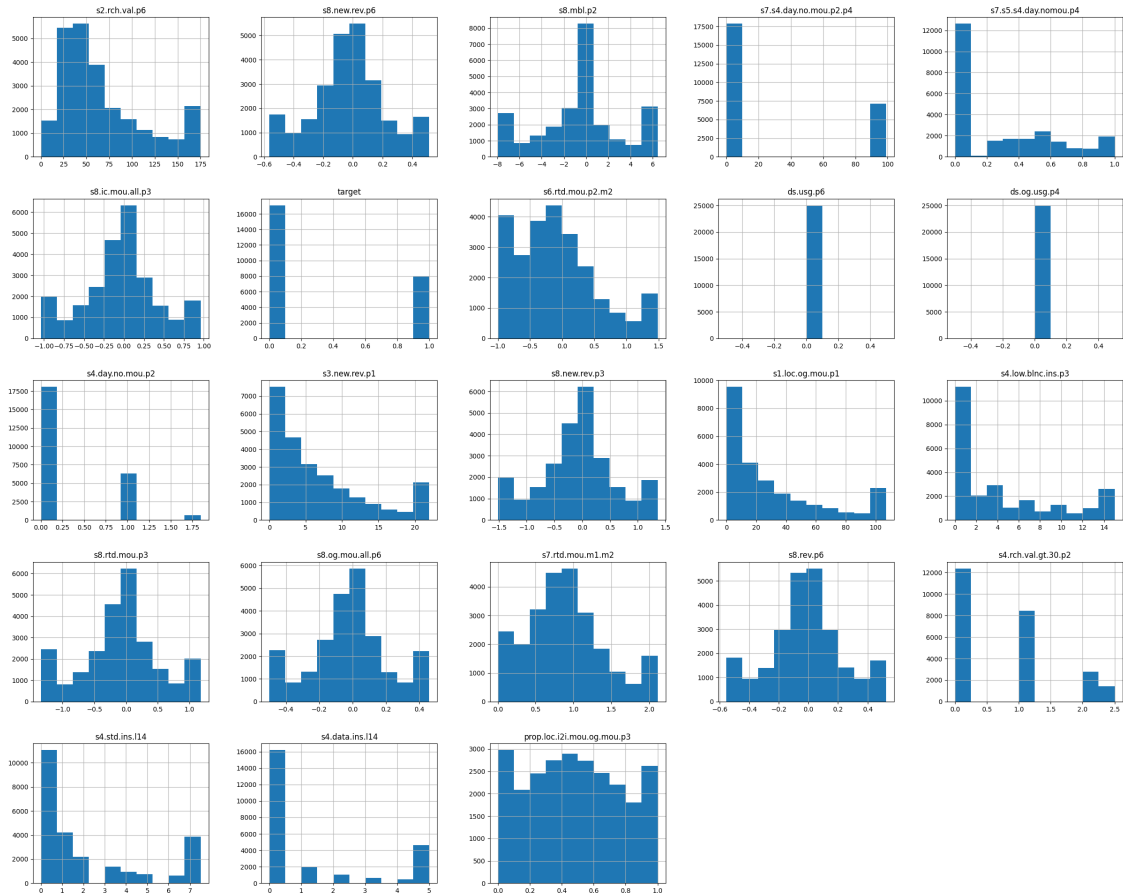
	prop.loc.i2i.mou.og.mou.p3
count	25000.000000
mean	0.483975
std	0.292349
min	0.000000
25%	0.251304
50%	0.477621
75%	0.716538
max	1.000000

[8 rows x 23 columns]

Final dataset shape: (25000, 23)

## 4.1 Visualize the Distribution

```
[35]: df.hist(figsize=(30,24))
plt.show()
```




```
[36]: df['target'].value_counts()
```

```
[36]: target
0    17083
1     7917
Name: count, dtype: int64
```

```
[37]: import plotly.graph_objects as go
fig=go.Figure(data=[go.Pie(labels=['Retained (0)', 'Exited_
↳(1)'],values=df['target'].value_counts())])
fig.update_layout(width=500, height=400)
fig.show()
```

```
[38]: s=df.select_dtypes(include=["integer","float"]).corr()
plt.figure(figsize=(15,15))
sns.heatmap(s,annot=True)
```

```
[38]: <Axes: >
```



## 4.2 Dropping the Irrelevant Columns

```
[39]: df = df.drop(['ds.og.usg.p4', 'ds.usg.p6'], axis=1)
      #df = df.drop(['ds.og.usg.p4', 's4.day.no.mou.p2', 'ds.usg.p6'], axis=1)
      #df = df.drop(['ds.usg.p6'], axis=1)
```

```
[40]: df.info()
```

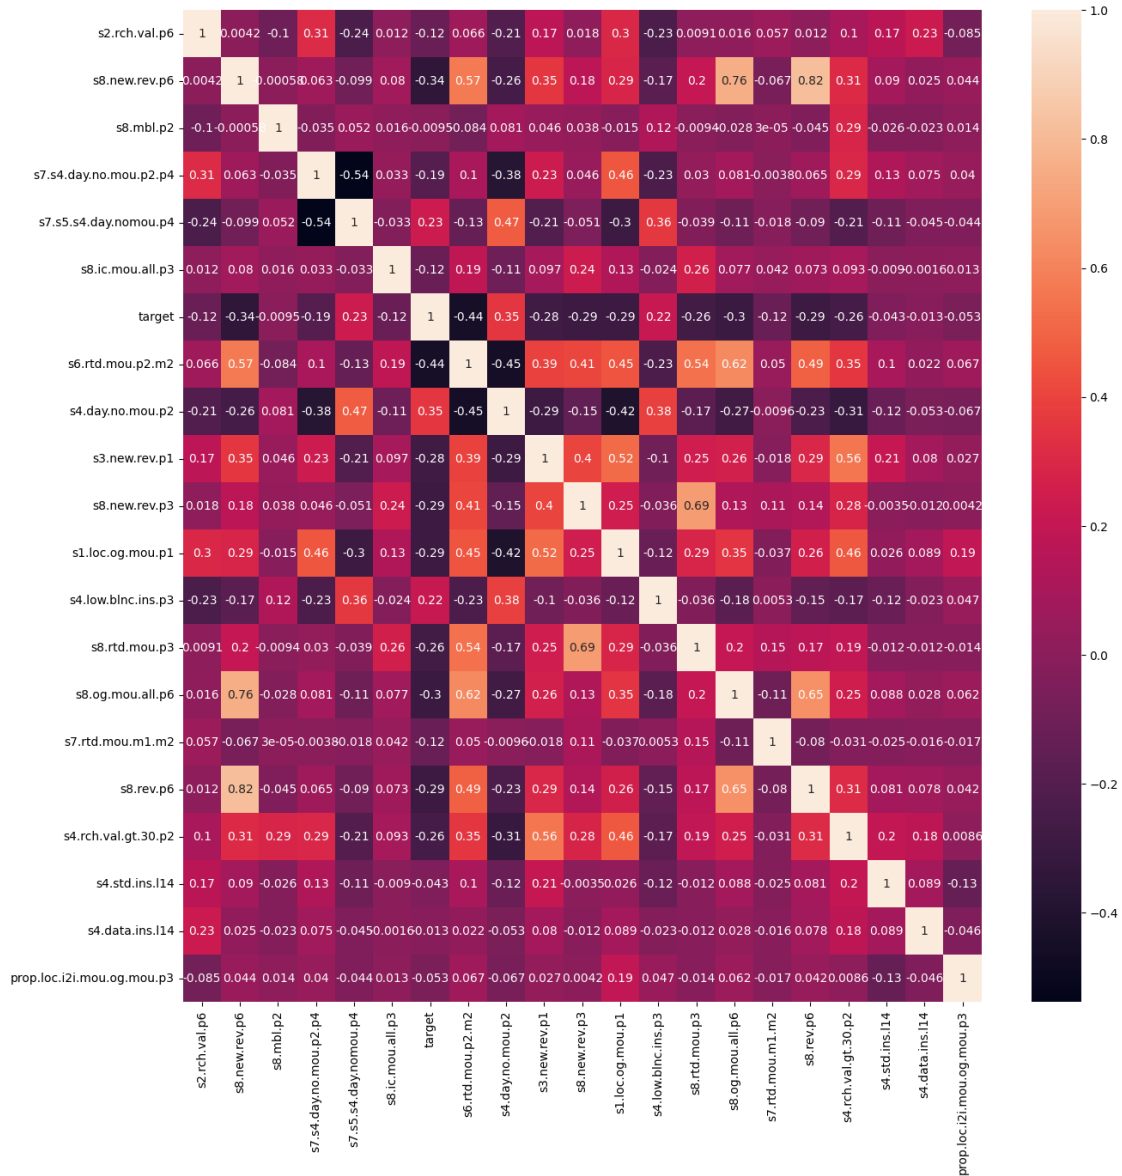
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25000 entries, 0 to 24999
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	s2.rch.val.p6	25000 non-null	float64
1	s8.new.rev.p6	25000 non-null	float64
2	s8.mbl.p2	25000 non-null	float64
3	s7.s4.day.no.mou.p2.p4	25000 non-null	float64
4	s7.s5.s4.day.nomou.p4	25000 non-null	float64
5	s8.ic.mou.all.p3	25000 non-null	float64
6	target	25000 non-null	int64
7	s6.rtd.mou.p2.m2	25000 non-null	float64
8	s4.day.no.mou.p2	25000 non-null	float64
9	s3.new.rev.p1	25000 non-null	float64
10	s8.new.rev.p3	25000 non-null	float64
11	s1.loc.og.mou.p1	25000 non-null	float64
12	s4.low.blnc.ins.p3	25000 non-null	int64
13	s8.rtd.mou.p3	25000 non-null	float64
14	s8.og.mou.all.p6	25000 non-null	float64
15	s7.rtd.mou.m1.m2	25000 non-null	float64
16	s8.rev.p6	25000 non-null	float64
17	s4.rch.val.gt.30.p2	25000 non-null	float64
18	s4.std.ins.l14	25000 non-null	float64
19	s4.data.ins.l14	25000 non-null	int64
20	prop.loc.i2i.mou.og.mou.p3	25000 non-null	float64

dtypes: float64(18), int64(3)  
memory usage: 4.0 MB

```
[41]: s=df.select_dtypes(include=["integer","float"]).corr()
plt.figure(figsize=(15,15))
sns.heatmap(s,annot=True)
```

```
[41]: <Axes: >
```



## 5 Feature Selection and Prepare the Data

```
[42]: X = df.drop('target', axis='columns')
      y = df['target']
```

```
[43]: X
```

```
[43]:      s2.rch.val.p6  s8.new.rev.p6  s8.mbl.p2  s7.s4.day.no.mou.p2.p4  \
0              39.29          -0.170          -0.72              1.000000
1              21.67          -0.320          -0.08              0.500000
```

2	30.00	-0.050	-0.09	0.384615
3	50.00	-0.180	1.83	0.416667
4	22.50	0.010	-0.04	0.222222
...	...	...	...	...
24995	26.67	-0.160	0.76	0.250000
24996	27.88	0.190	0.37	0.454545
24997	10.00	-0.030	-0.79	0.083333
24998	42.92	0.515	-1.09	99.000000
24999	53.50	-0.480	0.00	0.400000

	s7.s5.s4.day.nomou.p4	s8.ic.mou.all.p3	s6.rtd.mou.p2.m2	\
0	0.666667	-0.73	-0.71	
1	0.583333	0.00	-0.96	
2	0.384615	-1.03	-0.98	
3	0.250000	-0.43	-0.92	
4	0.777778	-1.04	-0.98	
...	...	...	...	
24995	0.375000	0.68	-0.58	
24996	0.636364	0.04	-0.07	
24997	0.333333	0.27	-0.12	
24998	0.000000	-0.42	0.25	
24999	0.333333	0.54	-1.00	

	s4.day.no.mou.p2	s3.new.rev.p1	s8.new.rev.p3	s1.loc.og.mou.p1	\
0	1.000000	0.22	-0.90	2.383200	
1	1.000000	0.38	-0.14	0.650000	
2	1.844649	0.11	-0.45	0.183300	
3	1.844649	0.49	-0.02	0.816500	
4	1.000000	0.10	-0.67	0.166600	
...	...	...	...	...	
24995	1.000000	0.76	0.26	3.499900	
24996	1.000000	14.26	0.82	17.116500	
24997	0.000000	8.39	0.38	3.299800	
24998	0.000000	15.16	-1.52	107.023513	
24999	1.000000	0.00	0.00	0.000000	

	s4.low.blnc.ins.p3	s8.rtd.mou.p3	s8.og.mou.all.p6	s7.rtd.mou.m1.m2	\
0	7	-0.500	-0.11	0.240533	
1	13	-0.110	-0.13	0.459725	
2	10	-0.390	-0.12	0.111785	
3	11	-0.020	-0.14	1.920826	
4	0	-0.630	-0.02	1.728186	
...	...	...	...	...	
24995	10	-0.130	-0.16	1.423358	
24996	10	0.220	0.08	0.688912	
24997	3	0.250	-0.03	1.223699	
24998	0	-1.345	0.46	0.579099	



```
24999          15          0.000          -0.50          1.423424
```

```

      s8.rev.p6  s4.rch.val.gt.30.p2  s4.std.ins.l14  s4.data.ins.l14  \
0          -0.120          0.0          0.0          0
1          -0.220          0.0          0.0          0
2          -0.070          0.0          1.0          0
3          -0.210          0.0          2.0          0
4           0.010          0.0          1.0          2
...          ...          ...          ...          ...
24995        -0.140          0.0          0.0          0
24996         0.050          0.0          0.0          0
24997        -0.030          0.0          1.0          1
24998         0.525          2.5          1.0          3
24999        -0.460          0.0          0.0          0

```

```

      prop.loc.i2i.mou.og.mou.p3
0          0.609456
1          1.000000
2          0.699592
3          0.086617
4          0.683105
...          ...
24995        0.328027
24996        0.288006
24997        0.235918
24998        0.952962
24999        1.000000

```

```
[25000 rows x 20 columns]
```

```
[44]: y
```

```

[44]: 0          1
      1          1
      2          1
      3          0
      4          0
      ..
24995        1
24996        0
24997        0
24998        0
24999        1
Name: target, Length: 25000, dtype: int64

```

```
[45]: X.shape
```

```
[45]: (25000, 20)
```

```
[46]: y.shape
```

```
[46]: (25000,)
```

## 5.1 Splitting the dataset

```
[77]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (20000, 20)
```

```
X_test shape: (5000, 20)
```

```
y_train shape: (20000,)
```


```
y_test shape: (5000,)
```

## 5.2 Check importance of Features in the Dataset

```
[78]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train,y_train.values.ravel())
```

```
[78]: RandomForestClassifier()
```

```
[79]: feat_scores = pd.DataFrame({"Fraction of variables affected" : rf.
    ↪feature_importances_},index=X.columns)
feat_scores = feat_scores.sort_values(by = "Fraction of variables affected")
feat_scores.plot(kind ="bar", figsize = (10,5))
sns.despine()
```



## 6 Train the Models

### 6.1 Random Forest Classifier

```
[80]: #from sklearn.ensemble import RandomForestClassifier
model_rf = RandomForestClassifier()
model_rf.fit(X_train,y_train)
```

```
[80]: RandomForestClassifier()
```

```
[81]: y_predict = model_rf.predict(X_test)
```


```
[82]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_predict))
```

	precision	recall	f1-score	support
0	0.82	0.90	0.86	3437
1	0.72	0.56	0.63	1563
accuracy			0.79	5000

macro avg	0.77	0.73	0.74	5000
weighted avg	0.79	0.79	0.79	5000

```
[83]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test,y_predict)
sns.heatmap(cm,annot=True)
```

[83]: <Axes: >




```
[84]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test,y_predict)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)

# Add labels and title
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')

# Add class labels
class_names = ['Negative (0)', 'Positive (1)']
tick_marks = np.arange(len(class_names))
```

```
plt.xticks(tick_marks + 0.5, class_names)
plt.yticks(tick_marks + 0.5, class_names)

# Show the plot
plt.tight_layout()
plt.show()
```



## 6.2 Linear Regression Model

```
[85]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import classification_report, confusion_matrix

      model_lr = LogisticRegression()
      model_lr.fit(X_train,y_train)
```

```
[85]: LogisticRegression()
```

```
[86]: y_predict = model_lr.predict(X_test)
```

```
[87]: print(classification_report(y_test,y_predict))
```


	precision	recall	f1-score	support
0	0.82	0.90	0.86	3437
1	0.72	0.56	0.63	1563
accuracy			0.80	5000
macro avg	0.77	0.73	0.74	5000
weighted avg	0.79	0.80	0.79	5000

```
[88]: cm = confusion_matrix(y_test,y_predict)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)

# Add labels and title
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')

# Add class labels
class_names = ['Negative (0)', 'Positive (1)']
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks + 0.5, class_names)
plt.yticks(tick_marks + 0.5, class_names)

# Show the plot
plt.tight_layout()
plt.show()
```



### 6.3 Support Vector Machine

```
[89]: from sklearn.calibration import CalibratedClassifierCV
      from sklearn.svm import LinearSVC

      model_svc = LinearSVC()
      model_svc.fit(X_train,y_train)
```

[89]: LinearSVC()

```
[90]: y_predict = model_svc.predict(X_test)
      print(classification_report(y_test, y_predict))
```


	precision	recall	f1-score	support
0	0.78	0.95	0.86	3437
1	0.78	0.42	0.55	1563
accuracy			0.78	5000
macro avg	0.78	0.68	0.70	5000
weighted avg	0.78	0.78	0.76	5000

```
[91]: cm = confusion_matrix(y_test,y_predict)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)

# Add labels and title
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')

# Add class labels
class_names = ['Negative (0)', 'Positive (1)']
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks + 0.5, class_names)
plt.yticks(tick_marks + 0.5, class_names)

# Show the plot
plt.tight_layout()
plt.show()
```





## 6.4 K- Nearest Neighbours

```
[92]: from sklearn.neighbors import KNeighborsClassifier
      from sklearn.metrics import classification_report, confusion_matrix
      model_knn = KNeighborsClassifier()
      model_knn.fit(X_train,y_train)
```

```
[92]: KNeighborsClassifier()
```

```
[93]: y_predict = model_knn.predict(X_test)
      print(classification_report(y_test, y_predict))
```


	precision	recall	f1-score	support
0	0.78	0.86	0.82	3437
1	0.61	0.48	0.53	1563
accuracy			0.74	5000
macro avg	0.69	0.67	0.68	5000
weighted avg	0.73	0.74	0.73	5000

```
[94]: cm = confusion_matrix(y_test,y_predict)
      sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)

      # Add labels and title
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
      plt.title('Confusion Matrix')

      # Add class labels
      class_names = ['Negative (0)', 'Positive (1)']
      tick_marks = np.arange(len(class_names))
      plt.xticks(tick_marks + 0.5, class_names)
      plt.yticks(tick_marks + 0.5, class_names)

      # Show the plot
      plt.tight_layout()
      plt.show()
```



## 6.5 Naive Bayer Model

```
[95]: from sklearn.naive_bayes import GaussianNB
      model_nb = GaussianNB()
      model_nb.fit(X_train,y_train)
```

```
[95]: GaussianNB()
```

```
[96]: y_predict = model_nb.predict(X_test)
```

```
[97]: print(classification_report(y_test, y_predict))
```


	precision	recall	f1-score	support
0	0.85	0.77	0.81	3437
1	0.58	0.71	0.64	1563
accuracy			0.75	5000
macro avg	0.72	0.74	0.73	5000
weighted avg	0.77	0.75	0.76	5000

```
[98]: cm = confusion_matrix(y_test,y_predict)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False)

# Add labels and title
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')

# Add class labels
class_names = ['Negative (0)', 'Positive (1)']
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks + 0.5, class_names)
plt.yticks(tick_marks + 0.5, class_names)

# Show the plot
plt.tight_layout()
plt.show()
```



## 7 ROC Curve

```
[99]: model_lr.predict_proba(X_test)
```

```
[99]: array([[0.99576888, 0.00423112],  
          [0.51117157, 0.48882843],  
          [0.80720218, 0.19279782],  
          ...,  
          [0.75609067, 0.24390933],  
          [0.83541985, 0.16458015],  
          [0.57106124, 0.42893876]])
```

```
[100]: model_lr.predict_proba(X_test)[: ,1]
```

```
[100]: array([0.00423112, 0.48882843, 0.19279782, ..., 0.24390933, 0.16458015,  
          0.42893876])
```

```
[101]: y_test
```

```
[101]: 6868      0  
      24016    0  
      9668    0  
     13640    0  
     14018    0  
      ...  
     8670    0  
     11839    0  
     4013    0  
     21147    0  
     695     0  
      Name: target, Length: 5000, dtype: int64
```

```
[111]: from sklearn.metrics import roc_curve  
fpr1,tpr1,thresh1 = roc_curve(y_test, model_rf.predict_proba(X_test)[:  
    ↪,1],pos_label=1)  
fpr2,tpr2,thresh2 = roc_curve(y_test, model_lr.predict_proba(X_test)[:  
    ↪,1],pos_label=1)  
fpr3,tpr3,thresh3 = roc_curve(y_test, model_svc.  
    ↪decision_function(X_test),pos_label=1)  
fpr4,tpr4,thresh4 = roc_curve(y_test, model_knn.predict_proba(X_test)[:  
    ↪,1],pos_label=1)  
fpr5,tpr5,thresh5 = roc_curve(y_test, model_nb.predict_proba(X_test)[:  
    ↪,1],pos_label=1)
```

```
[112]: from sklearn.metrics import roc_auc_score  
roc_auc_score1 = roc_auc_score(y_test, model_rf.predict_proba(X_test)[: ,1])  
roc_auc_score2 = roc_auc_score(y_test, model_lr.predict_proba(X_test)[: ,1])
```

```

roc_auc_score3 = roc_auc_score(y_test, model_svc.decision_function(X_test))
roc_auc_score4 = roc_auc_score(y_test, model_knn.predict_proba(X_test)[: ,1])
roc_auc_score5 = roc_auc_score(y_test, model_nb.predict_proba(X_test)[: ,1])

print("Random Forest:", roc_auc_score1)
print("Logistic Regression", roc_auc_score2)
print("Support Vector Machine", roc_auc_score3)
print("K-Nearest Neighbours", roc_auc_score4)
print("Naive Bayes:", roc_auc_score5)

```

```

Random Forest: 0.8373767761206145
Logistic Regression 0.8340912403521127
Support Vector Machine 0.7929215226047653
K-Nearest Neighbours 0.7074472392285152
Naive Bayes: 0.8148504727541593

```

```


[113]: plt.plot(fpr1,tpr1,linestyle='--', color='orange', label='Random Forest')
plt.plot(fpr2,tpr2,linestyle='--', color='green', label='Linear Regression')
plt.plot(fpr3,tpr3,linestyle='--', color='yellow', label='Support Vector
Machine')
plt.plot(fpr4,tpr4,linestyle='--', color='red', label='K-Nearest Neighbours')
plt.plot(fpr5,tpr5,linestyle='--', color='blue', label='Naive Bayes')

```

```

[113]: [<matplotlib.lines.Line2D at 0x7963d4b13ee0>]

```



## 8 Conclusion

As per the Graph and ROC AUC Score, Random Forest Model performed Well.